BCNF, Memory Hierarchy & File Structures

Exercise: BCNF

- BCNF: Boyce-Codd Normal Form: Boyce-Codd范式
- R is in BCNF if and only if for each FD: X → {A} in F+, A ∈ X (trivial FD), or X is a superkey for R
- In other words: 对于每个不包含无关(无用)属性的 FD,每个 FD 的 左侧 都是候选键。

Exercise: BCNF (cont)

$$R = (A, B, C, D).$$

 $F = \{C \rightarrow D, C \rightarrow A, B \rightarrow C\}.$
Decompose R into a set of BCNF relations.

- $S=\{R\}$. When we find a schema R with BCNF violation $X \to Y$ we:
 - 1. 把R从S中删除
 - 2. 新建一个schema加入S,其属性集为 R Y
 - 3. 新建一个schema加入S, 其属性集为 X和Y的并集

Exercise: BCNF (cont)

$$R = (A, B, C, D).$$

 $F = \{C \rightarrow D, C \rightarrow A, B \rightarrow C\}.$

Decompose R into a set of BCNF relations.

- List all candidate keys of R.
 {B}
- $C \rightarrow D$, $C \rightarrow A$ violates BCNF. Because C is not a super key for R.

Take C
$$\rightarrow$$
 D: decompose R to $R_1 = \{A, \underline{B}, C\}, R_2 = \{\underline{C}, D\}.$

• $C \rightarrow A$ in R_1 violates BCNF. C is not a super key for R_1 .

Decompose R₁ to

$$R_{11} = \{ \underline{B}, C \} R_{12} = \{ \underline{C}, A \}.$$

Final decomposition:

$$R_2 = \{\underline{C}, D\}, R_{11} = \{\underline{B}, C\}, R_{12} = \{\underline{C}, A\}.$$

Memory Hierarchy

| | Cost (money) | Access (time) | volatile | Capacity |
|-------------|------------------|---------------|----------|----------|
| Cache | Very expensive | Very Fast | Yes | KB -> MB |
| Main memory | Expensive | Fast | Yes | MB -> GB |
| Disk \ | Inexpensive | Slow | No | GB -> TB |
| Tape | Very inexpensive | Very slow | No | TB -> PB |
| | | | | |

Database

ref: <u>numbers</u>

Accessing Disk

- Access unit: disk blocks or pages (typically a block is 1K-16Kbytes)
- Physical movement, random access
- Time to access (read/write) a disk block seek time + rotational delay + transfer time

Database<->a collection of files

File<->a sequence of records

Record<->a sequence of fields(属性、域值)

Storing records in files

- † Fixed-Length Records
 - Free Lists
- Variable-Length Records
 - Byte String Representation
 - Reserved space
 - Pointer Method
 - Slotted Page Structure

Fixed-Length Records

Fixed-Length Records

Record access is simple: n*(i-1)

Problem: what if deletion?

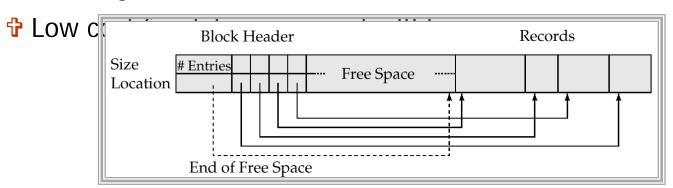
Free list

| header | | | | _ | |
|----------|-------|------------|-----|---|----------|
| record 0 | A-102 | Perryridge | 400 | |) |
| record 1 | | | | _ | \prec |
| record 2 | A-215 | Mianus | 700 | |) |
| record 3 | A-101 | Downtown | 500 | | |
| record 4 | | | | 1 | |
| record 5 | A-201 | Perryridge | 900 | | |
| record 6 | | | | _ | 4 |
| record 7 | A-110 | Downtown | 600 | | _ |
| record 8 | A-218 | Perryridge | 700 | | |

- Variable-Length Records
 - Byte String Representation : Attach an end-of-record (⊥)
 control character to the end of each record
 - Problem: deletion & record growth
 - Reserved space: use fixed-length records of a known maximum length
 - Problem: waste storage space
 - Pointer Method: Anchor block + Overflow block

| anchor block | Perryridge | A-102 | 400 | |
|--------------|------------|-------|-----|---|
| | Round Hill | A-305 | 350 | |
| | Mianus | A-215 | 700 | |
| | Downtown | A-101 | 500 | |
| | Redwood | A-222 | 700 | |
| | Brighton | A-217 | 750 | X |
| | | | | |
| overflow | V | A-201 | 900 | |
| block | | A-218 | 700 | |
| | | A-110 | 600 | |

- Variable-Length Records
 - Byte String Representation : Attach an end-of-record (⊥)
 control character to the end of each record
 - Problem: deletion & record growth
 - Reserved space: use fixed-length records of a known maximum length
 - Problem: waste storage space
 - Pointer Method: Anchor block + Overflow block
 - Slotted Page Structure:



Database<->a collection of files

File<->a sequence of records

Record<->a sequence of fields

- Heap
 - record can be placed anywhere in the file where there is space
- Sequential
 - store records in sequential order
- Hashing
 - a hash function computed on some attribute of each record

- Assume that a school keeps a file with the records of its students: Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes), dept-id is the department(院系) id where a student belongs to.
- There exist 10,000 student records and 50 departments. A page is 128 bytes. The data file is sorted sequentially on sid.
- **Q**1: What's the size of a record?
- Q2: How many pages do we need to store these student records?
- Q3: Given this data file, what's the cost to find a particular student by sid?

- There exist 10,000 student records and 50 departments. A page is 128 bytes and a pointer is 4 bytes. The data file is sorted sequentially on sid.
- **\dagger** What's the size of a record?

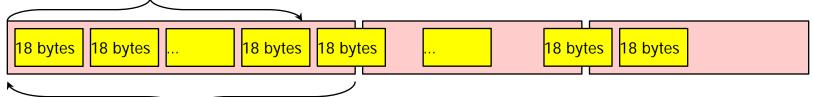
$$4 + 10 + 4 = 18$$

- Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes)
- There exist 10,000 student records and 50 departments. A page is 128 bytes and a pointer is 4 bytes. The data file is sorted sequentially on sid.
- How many pages do we need to stor these student records?

```
(#records * size of record) / size of page = 10,000 * 18 / 128 = 1406.25 \approx 1407
```

7 records, of size 18*7=126 bytes

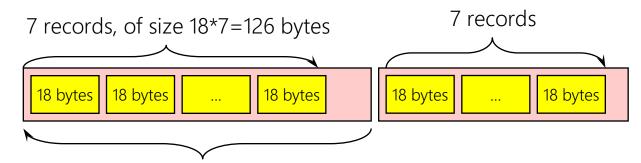
不允许记录跨越页面存储



Page size: 128 bytes

- Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes)
- There exist 10,000 student records and 50 departments. A page is 128 bytes and a pointer is 4 bytes. The data file is sorted sequentially on sid.
- How many pages do we need to store these student records?

[(# of records /
$$\lfloor$$
 size of page / size of record \rfloor)] = $\lceil (10,000 / \lfloor 128 / 18 \rfloor) \rceil = 1429$



Page size: 128 bytes

- Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes)
- There exist 10,000 student records and 50 departments. A page is 128 bytes and a pointer is 4 bytes. The data file is sorted sequentially on sid.
- Q3: Given this data file, what's the cost to find a particular student by sid?

of pages required: 1429

Binary search on sid:

$$\lceil \log_2(\# \text{ of pages}) \rceil = \lceil \log_2 1429 \rceil = 11$$