

Tutorial 7



Indexing, B⁺-tree, and Hashing

Part 1: Indexing

- ❖ Speed up access to desired data.
- ❖ **Search Key** - attribute to set of attributes used to look up records in a file.

Two kinds of index

- ❖ Ordered index
- ❖ Hash Index

Index

- ❖ **Primary index (clustering index):** in a sequentially ordered file, the index whose search key specifies the sequential order of the file.
- ❖ **Secondary index (non-clustering index):** an index whose search key specifies an order different from the sequential order of the file.
- ❖ **Dense index:** one index, one record
- ❖ **Sparse Index:** one index, many records; only applicable to clustering index
- ❖ **Single level Index:** index for data
- ❖ **Multilevel Index:** index of index

Exercise 1

- ❖ 假设一所学校保存了一份包含学生记录的档案: Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes), dept-id is the department(院系) id where a student belongs to.
- ❖ There exist 10,000 student records and 50 departments. A page is 128 bytes and a pointer is 4 bytes. The data file is sorted sequentially on sid.
- ❖ Q1: Given the data file only, what's the cost of finding students in a particular department (e.g., CSE)?
- ❖ Q2: How to improve?

Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes)

Exercise 1

- ❖ Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes)
- ❖ There exist 10,000 student records and 50 departments. A page is 128 bytes and a pointer is 4 bytes. The data file is sorted sequentially on sid.
- ❖ Q1: Given the data file only, what's the cost of finding students in a particular department (e.g., CSE)?

Pages required to store data file is 1429.

Records are sorted on sid, instead of dept-id.

Given the data file, the only way is to sequentially scan it.

The cost is 1429.

Exercise 1

- ❖ Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes)
- ❖ There exist 10,000 student records and 50 departments. A page is 128 bytes and a pointer is 4 bytes. The data file is sorted sequentially on sid.
- ❖ Q2: How to improve?

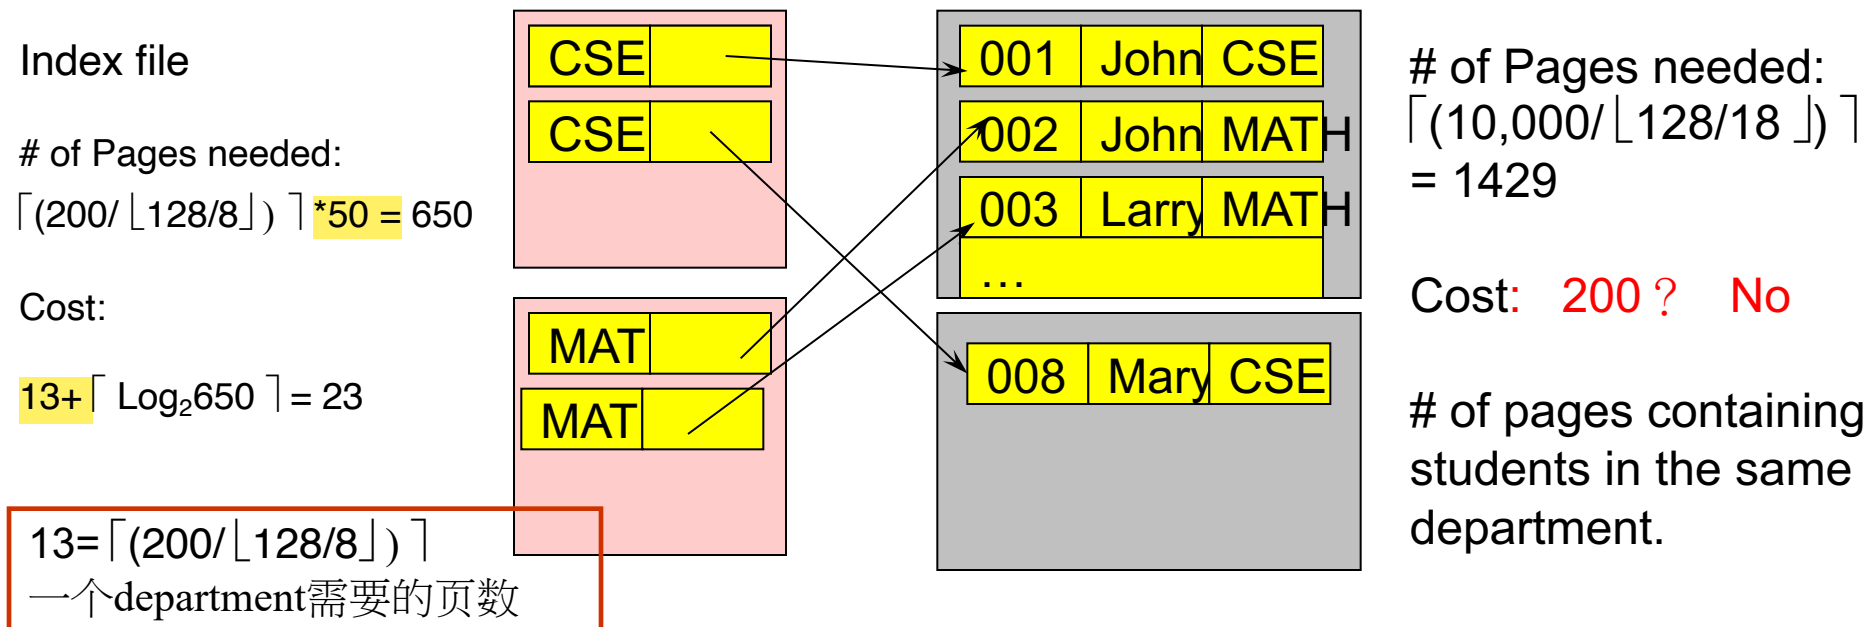
Build index on dept-id.

The entry for an index: (dept-id, pointer), each is of size 8 bytes.

Exercise 1


- ❖ Student (sid:4 bytes, sname: 10 bytes, dept-id: 4 bytes)
- ❖ There exist 10,000 student records and 50 departments. A page is 128 bytes and a pointer is 4 bytes. The data file is sorted sequentially on sid.
- ❖ 为了简单起见，我们假设每个系正好有 200 名学生。并且在索引页面中，一个页面只能包含同一系学生的指针。

❖ Q2-1: Build ordered index.



Exercise 1

❖ Q2-1: Build ordered index.



Primary index or secondary index?


Secondary index.

Dense index or sparse index?

Dense index.

Secondary index must be dense index.

Exercise 1




Primary index or secondary index?

Primary index.

Dense index or sparse index?

Dense index.

Exercise 1



Primary index or secondary index?

Primary index.

Dense index or sparse index?

Sparse index.

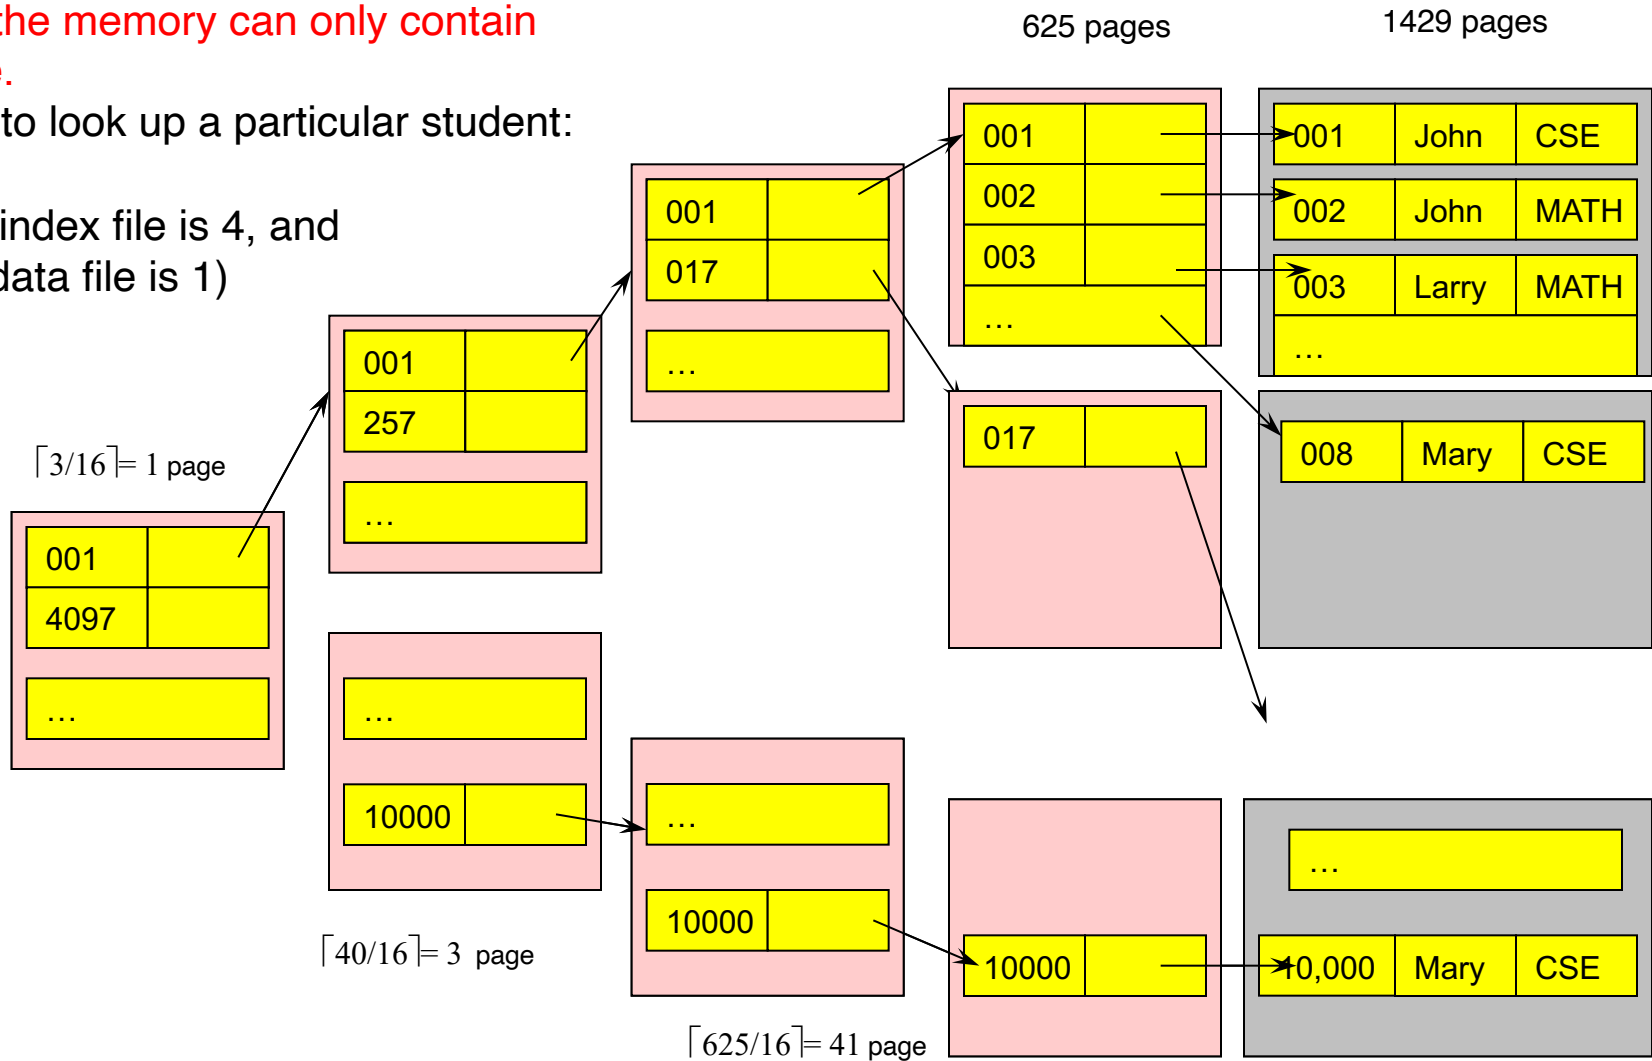
Exercise 1

Assume the memory can only contain one page.

The cost to look up a particular student:

$$4 + 1 = 5$$

(Cost on index file is 4, and cost on data file is 1)



一个page可以装 $128/8 = 16$ 个索引

Hash Index

- ❖ Hashing can be used not only for data file organization, but also for index-structure creation.
- ❖ A **hash index** organizes the search keys, with their associated record pointers, into a hash file structure.
- ❖ Strictly speaking, **hash indices are always secondary indices.**

Hashing: Hash Functions

❖ Worst case

- ❖ Hash function maps all search-key values to the same bucket
- ❖ Access time proportional to the number of search-key values
- ❖ Data: 2, 7, 12, 17, 22...
- ❖ Hash function $F(x) = x \bmod 5$

| Bucket 0 | Bucket 1 | Bucket 2 | Bucket 3 | Bucket 4 |
|----------|----------|----------|----------|----------|
| | | 2 | | |
| | | 7 | | |
| | | 12 | | |
| | | 17 | | |
| | | 22 | | |

Hashing: Hash Functions

❖ Ideal hash function

- ❖ Random

- ❖ Each bucket receives same number of records

- ❖ Data: 0, 1, 2, 3, 4, 5, 6, 7, 8, 9...

- ❖ Hash function $F(x) = x \bmod 5$

| Bucket 0 | Bucket 1 | Bucket 2 | Bucket 3 | Bucket 4 |
|----------|----------|----------|----------|----------|
| 0 | 1 | 2 | 3 | 4 |
| 5 | 6 | 7 | 8 | 9 |
| | | | | |
| | | | | |
| | | | | |

Handling of Bucket Overflows


- ❖ Although the probability of bucket overflow can be reduced, it cannot be eliminated; it is handled by using *overflow buckets*.
- ❖ **Overflow chaining** — 溢出桶在链表中链接在一起。长链会降低性能，因为查询必须读取链中的所有存储桶。

Part 2 B⁺-tree and Hashing

- ❖ B⁺-tree
- ❖ Dynamic Hashing

Review: B⁺-tree

- ❖ Balanced tree
- ❖ Use pointer, no need for sequential storage
- ❖ 树的深度是对数 (\log) 级别
- ❖ **Space overhead** (空间额外开销, 存储树索引)
- ❖ **Insertion and deletion overhead.** However, can be handled in logarithmic time. (插入删除的开销, 由维护树结构导致, 但可以控制在对数时间内)




Review: Extendible Hashing

- ❖ Allow the number of buckets to be modified dynamically
- ❖ Insert: If bucket is full, split it
- ❖ If necessary, double the directory
 - ❖ Before inserting, **local depth** of bucket = **global depth**
 - ❖ Insert causes **local depth** > **global depth**
- ❖ Removal of data entry makes bucket empty
 - ❖ **Merge "split image"**
- ❖ Each directory element points to same bucket as its split image
 - ❖ **Halve directory**

Exercise 2: B⁺-tree


❖ Given a B⁺-tree:




❖ Draw the tree after

1. Inserting 8
2. Deleting 2
3. Deleting 3

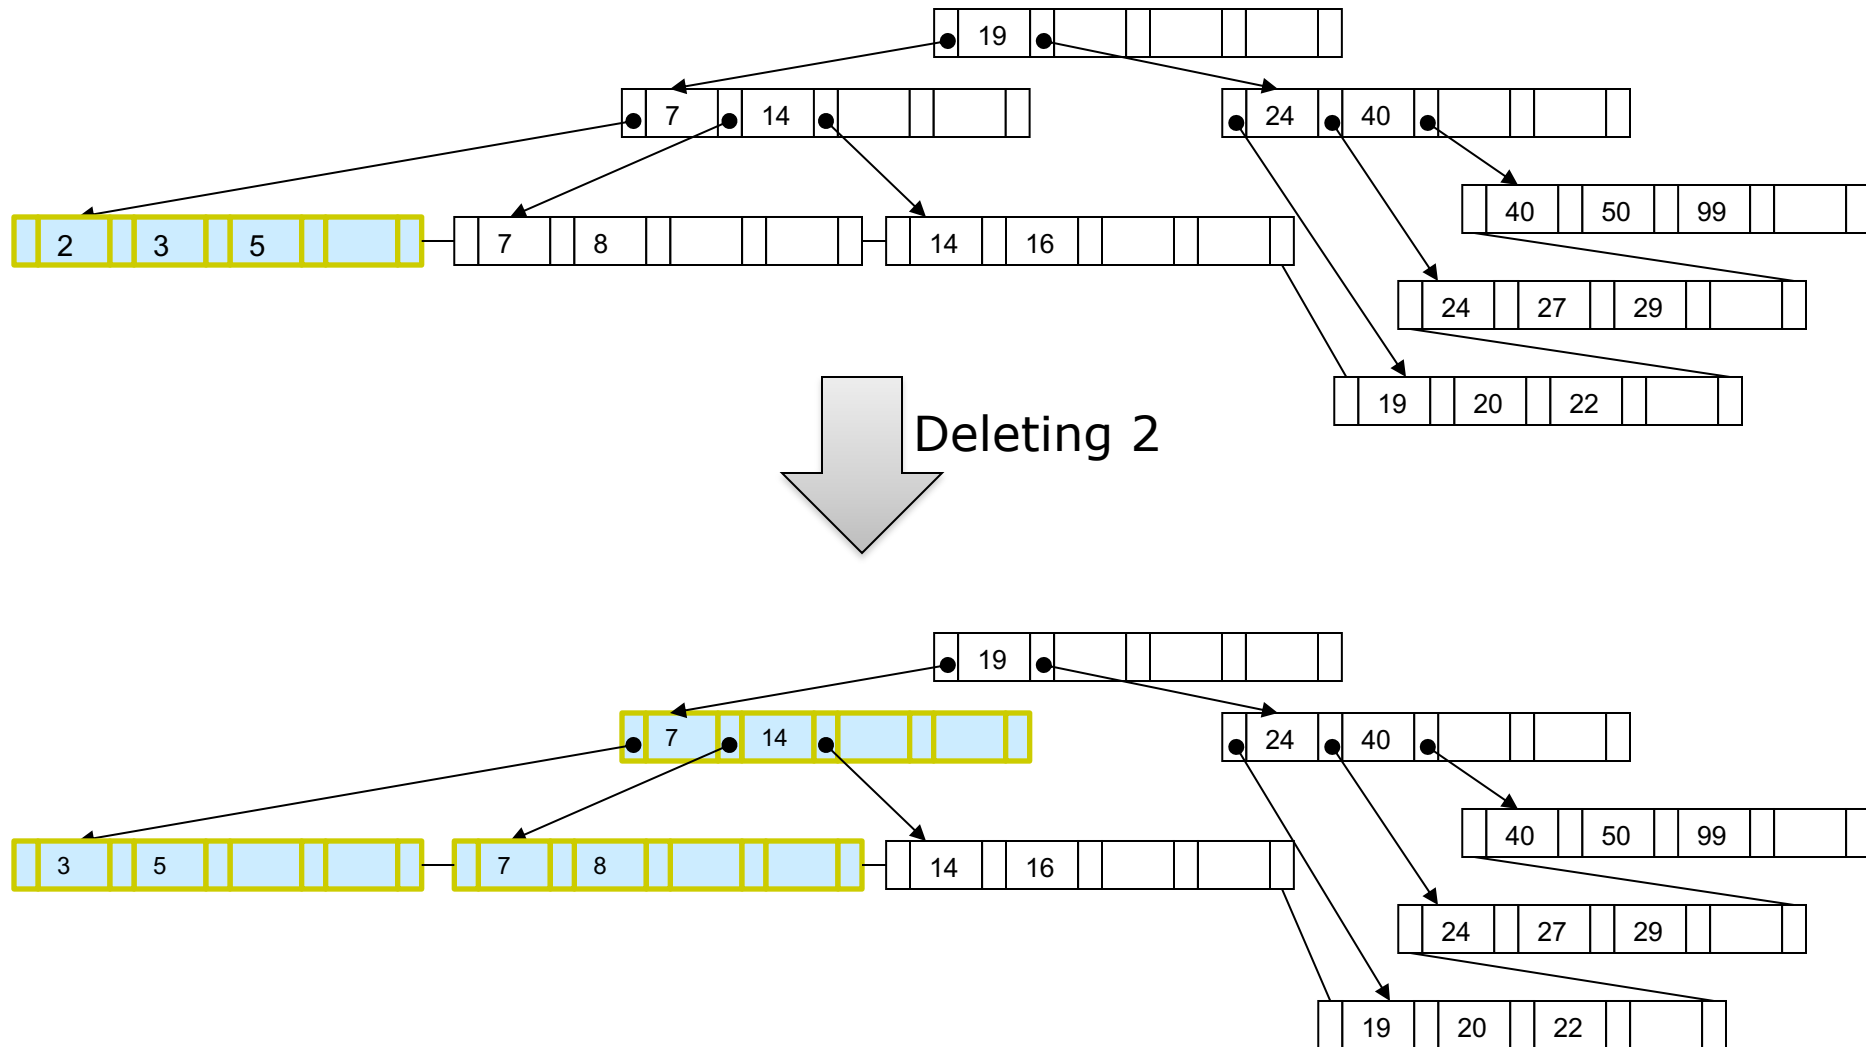
Exercise 2: B⁺-tree (1/3)



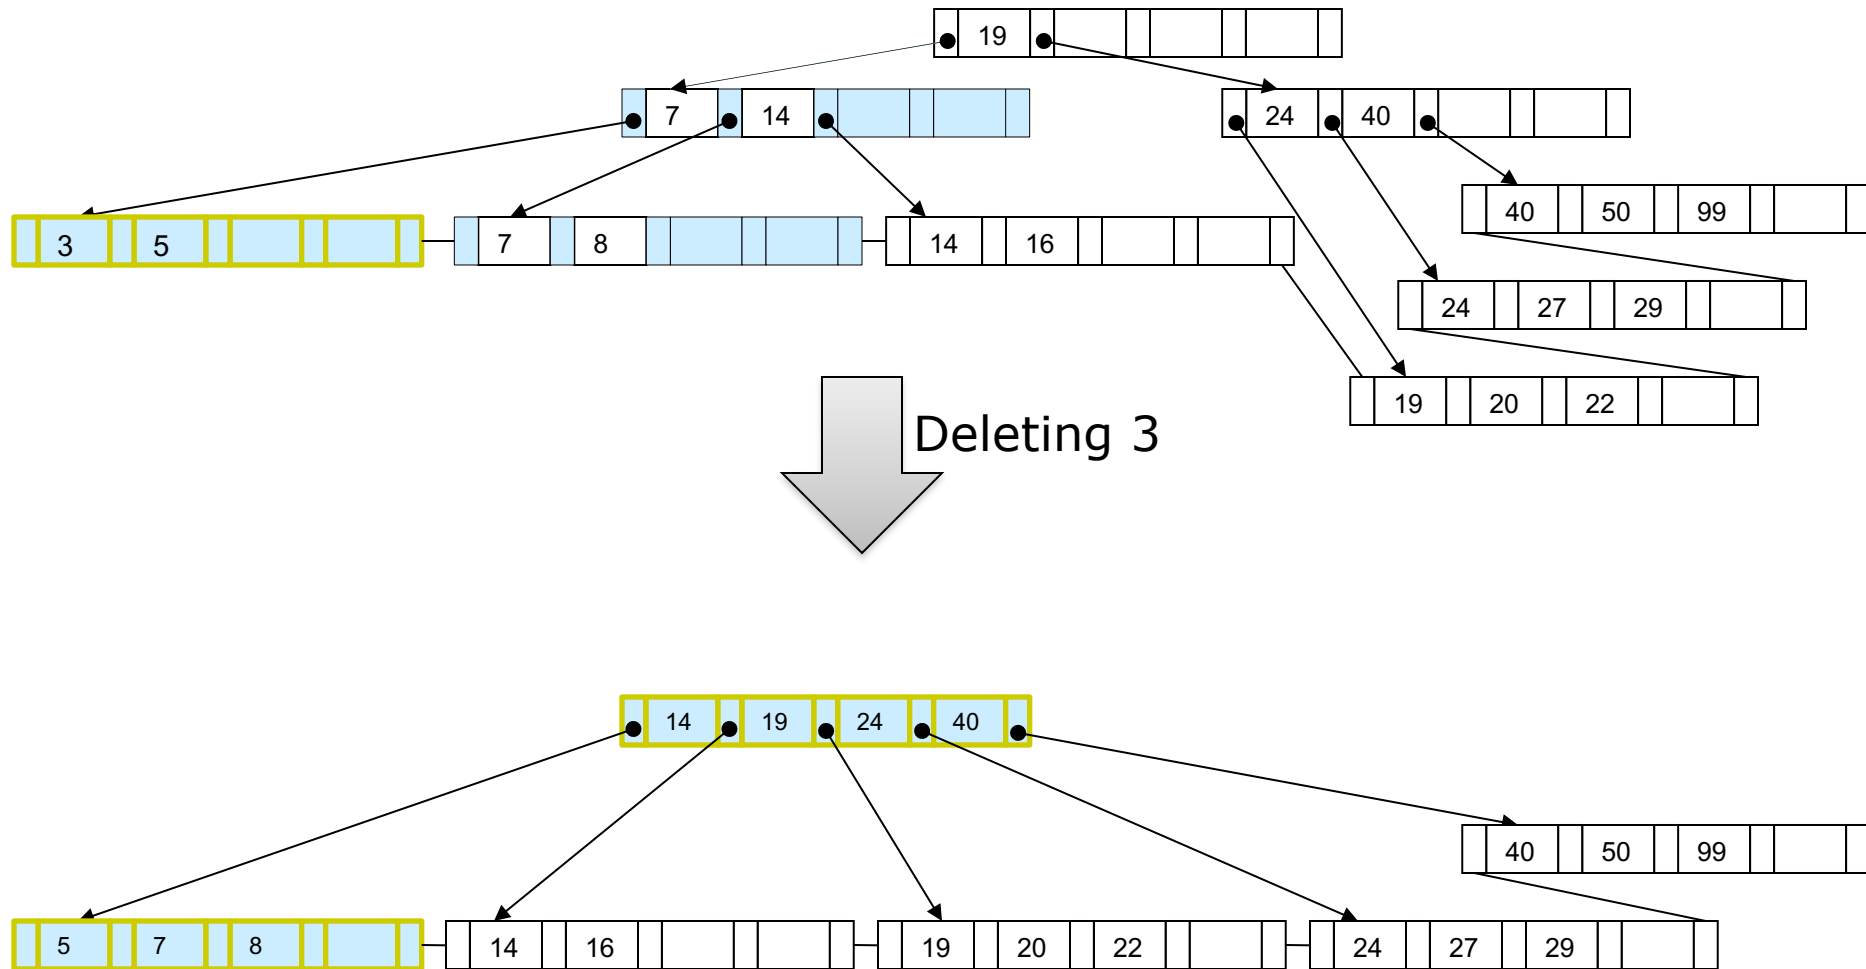
Inserting 8



Exercise 2: B⁺-tree (2/3)




Exercise 2: B⁺-tree (3/3)



Exercise 3: Extendible Hashing

❖ Given the following directory and buckets.



❖ Using extendible hashing, what will they be after:




1. Inserting 22
2. Inserting 3
3. Inserting 9

Exercise 3: Extendible Hashing (1/3)

❖ Inserting 22 (000101**10**)

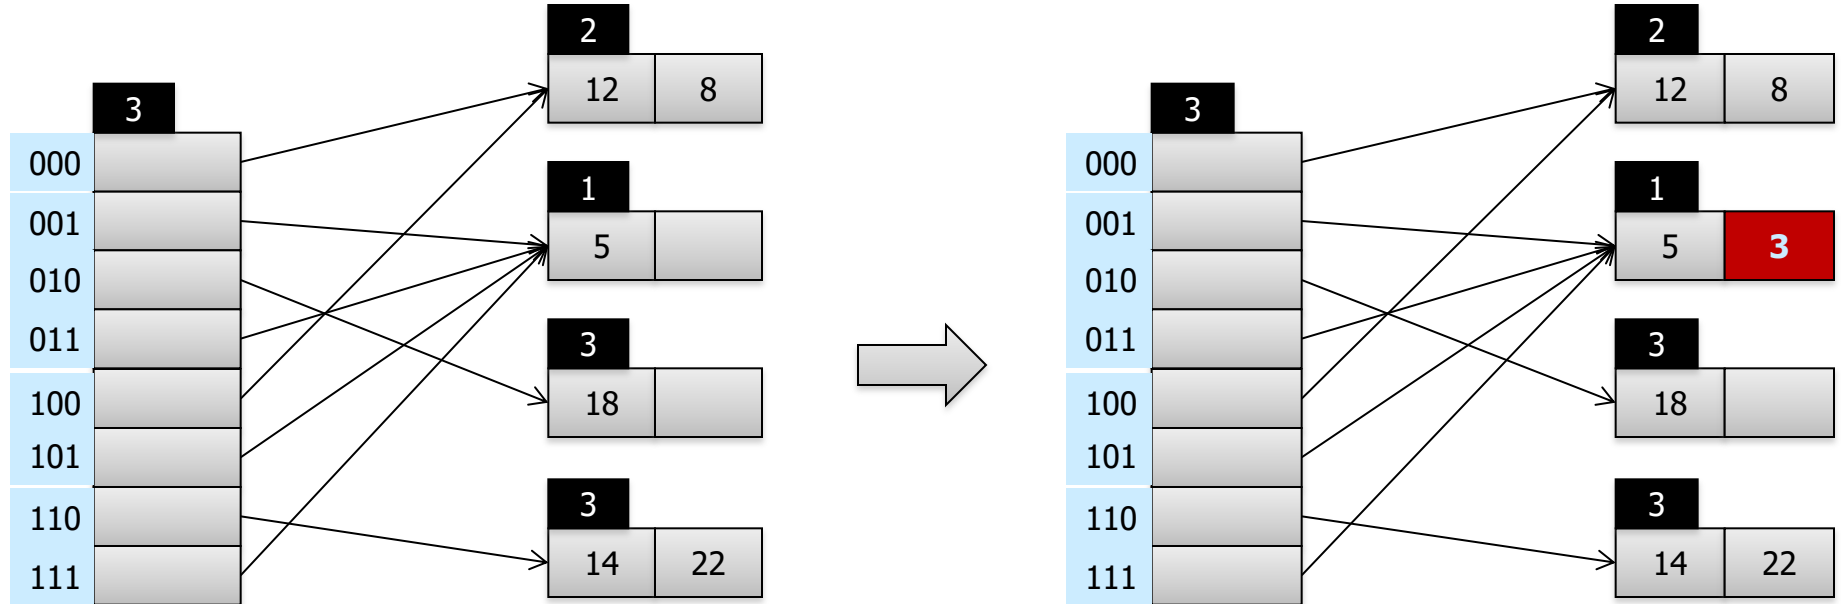
❖ 14 (0000 1**10**)

❖ 18 (0001 0**010**)



Exercise 3: Extendible Hashing (2/3)

❖ Inserting 3 (0000 00**11**)



Exercise 3: Extendible Hashing (3/3)

❖ Inserting 9 (0000 1001)

